

Part 1

1. MyGRUCell implementation

```
class MyGRUCell(nn.Module):
    def __init__(self, input_size, hidden_size):
        super(MyGRUCell, self).__init__()

        self.input_size = input_size
        self.hidden_size = hidden_size

        # -----
        # FILL THIS IN
        # -----
        ## Input linear layers
        self.Wiz = nn.Linear(input_size, hidden_size)
        self.Wir = nn.Linear(input_size, hidden_size)
        self.Wih = nn.Linear(input_size, hidden_size)

        ## Hidden linear layers
        self.Whz = nn.Linear(hidden_size, hidden_size)
        self.Whr = nn.Linear(hidden_size, hidden_size)
        self.Whh = nn.Linear(hidden_size, hidden_size)

    def forward(self, x, h_prev):
        """Forward pass of the GRU computation for one time step.

        Arguments
            x: batch_size x input_size
            h_prev: batch_size x hidden_size

        Returns:
            h_new: batch_size x hidden_size
        """

        # -----
        # FILL THIS IN
        # -----
        z = torch.sigmoid(self.Wiz(x) + self.Whz(h_prev))
        r = torch.sigmoid(self.Wir(x) + self.Whr(h_prev))
        g = torch.tanh(self.Wih(x) * self.Whh(h_prev))
        h_new = (1-z) * g + z * h_prev
        return h_new
```

2. When I tried to translate words starting with t, it does not perform well:

tea team tight -> eatay eaway ightedway

Also, when I tried to translate words ending with 'ing', the model also does not perform well:

shopping fighting running -> osingpay ightfay-indedway
uningway

Part 2

1. The equations are as follow:

$$\tilde{\alpha}_i^{(t)} = f(Q_t, K_i) = W_2(\text{ReLU}(W_1[Q_t K_i] + b1)) + b2$$

$$\alpha_i^{(t)} = \text{softmax}(\tilde{\alpha}_i^{(t)})$$

$$c_t = \sum_{i=1}^T \alpha_i^{(t)} K_i$$

2. My RNNAttentionDecoder implementation

```
def forward(self, inputs, annotations, hidden_init):
    """Forward pass of the attention-based decoder RNN.

    Arguments:
        inputs: Input token indexes across a batch for all the time step. (batch_size x decoder_seq_len)
        annotations: The encoder hidden states for each step of the input.
                    sequence. (batch_size x seq_len x hidden_size)
        hidden_init: The final hidden states from the encoder, across a batch. (batch_size x hidden_size)

    Returns:
        output: Un-normalized scores for each token in the vocabulary, across a batch for all the decoding time steps. (batch
        attentions: The stacked attention weights applied to the encoder annotations (batch_size x encoder_seq_len x decoder_
        """"

    batch_size, seq_len = inputs.size()
    embed = self.embedding(inputs) # batch_size x seq_len x hidden_size

    hiddens = []
    attentions = []
    h_prev = hidden_init
    for i in range(seq_len):
        # -----
        # FILL THIS IN - START
        # -----
        embed_current = embed[:,i,:] # Get the current time step, across the whole batch (batch size x hidden size)
        context, attention_weights = self.attention(embed_current, annotations, annotations) # batch_size x 1 x hidden_size
        embed_and_context = torch.cat([embed_current, torch.squeeze(context, dim=1)], dim=1) # batch_size x (2*hidden_size)
        h_prev = self.rnn(embed_and_context, h_prev) # batch_size x hidden_size
        # -----
        # FILL THIS IN - END
        # -----

        hiddens.append(h_prev)
        attentions.append(attention_weights)

    hiddens = torch.stack(hiddens, dim=1) # batch_size x seq_len x hidden_size
    attentions = torch.cat(attentions, dim=2) # batch_size x seq_len x seq_len

    output = self.out(hiddens) # batch_size x seq_len x vocab_size
    return output, attentions
```

Part 3

1. ScaledDotProduct implementation

```

class ScaledDotAttention(nn.Module):
    def __init__(self, hidden_size):
        super(ScaledDotAttention, self).__init__()

        self.hidden_size = hidden_size

        self.Q = nn.Linear(hidden_size, hidden_size)
        self.K = nn.Linear(hidden_size, hidden_size)
        self.V = nn.Linear(hidden_size, hidden_size)
        self.softmax = nn.Softmax(dim=1)
        self.scaling_factor = torch.rsqrt(torch.tensor(self.hidden_size, dtype= torch.float))

    def forward(self, queries, keys, values):
        """The forward pass of the scaled dot attention mechanism.

        Arguments:
            queries: The current decoder hidden state, 2D or 3D tensor. (batch_size x (k) x hidden_size)
            keys: The encoder hidden states for each step of the input sequence. (batch_size x seq_len x hidden_size)
            values: The encoder hidden states for each step of the input sequence. (batch_size x seq_len x hidden_size)

        Returns:
            context: weighted average of the values (batch_size x k x hidden_size)
            attention_weights: Normalized attention weights for each encoder hidden state. (batch_size x seq_len x k)

        The output must be a softmax weighting over the seq_len annotations.
        """

        # -----
        # FILL THIS IN
        # -----
        batch_size = keys.size(0)
        q = self.Q(queries)
        k = self.K(keys)
        v = self.V(values)
        unnormalized_attention = (k @ q.transpose(1,2))/ self.scaling_factor
        attention_weights = self.softmax(unnormalized_attention)
        context = attention_weights.transpose(1,2) @ v
        return context, attention_weights

```

CausalScaledDotProduct implementation

```

class CausalScaledDotAttention(nn.Module):
    def __init__(self, hidden_size):
        super(CausalScaledDotAttention, self).__init__()

        self.hidden_size = hidden_size
        self.neg_inf = torch.tensor(-1e7)

        self.Q = nn.Linear(hidden_size, hidden_size)
        self.K = nn.Linear(hidden_size, hidden_size)
        self.V = nn.Linear(hidden_size, hidden_size)
        self.softmax = nn.Softmax(dim=1)
        self.scaling_factor = torch.rsqrt(torch.tensor(self.hidden_size, dtype= torch.float))

    def forward(self, queries, keys, values):
        """The forward pass of the scaled dot attention mechanism.

        Arguments:
            queries: The current decoder hidden state, 2D or 3D tensor. (batch_size x (k) x hidden_size)
            keys: The encoder hidden states for each step of the input sequence. (batch_size x seq_len x hidden_size)
            values: The encoder hidden states for each step of the input sequence. (batch_size x seq_len x hidden_size)

        Returns:
            context: weighted average of the values (batch_size x k x hidden_size)
            attention_weights: Normalized attention weights for each encoder hidden state. (batch_size x seq_len x k)

        The output must be a softmax weighting over the seq_len annotations.
        """

        # -----
        # FILL THIS IN
        # -----
        batch_size = keys.size(0)
        q = self.Q(queries)
        k = self.K(keys)
        v = self.V(values)
        unnormalized_attention = (k @ q.transpose(1,2))/ self.scaling_factor
        mask = torch.tril(unnormalized_attention)
        mask[mask == 0] = self.neg_inf
        attention_weights = self.softmax(mask)
        context = attention_weights.transpose(1,2) @ v
        return context, attention_weights

```

TransformerEncoder's forward implementation

```
def forward(self, inputs):
    """Forward pass of the encoder RNN.

    Arguments:
        inputs: Input token indexes across a batch for all time steps in the sequence. (batch_size x seq_len)

    Returns:
        annotations: The hidden states computed at each step of the input sequence. (batch_size x seq_len x hidden_size)
        hidden: The final hidden state of the encoder, for each sequence in a batch. (batch_size x hidden_size)
    """

    batch_size, seq_len = inputs.size()
    # -----
    # FILL THIS IN - START
    # -----
    encoded = self.embedding(inputs) # batch_size x seq_len x hidden_size

    # Add positional embeddings from self.create_positional_encodings. (a'la https://arxiv.org/pdf/1706.03762.pdf, section 3.5)
    encoded += self.positional_encodings[:seq_len]

    annotations = encoded

    for i in range(self.num_layers):
        new_annotations, self_attention_weights = self.self_attentions[i](encoded, annotations, annotations) # batch_size x seq_len x hidden_size
        residual_annotations = annotations + new_annotations
        new_annotations = self.attention_mlp[i](residual_annotations)
        annotations = residual_annotations + new_annotations
    # -----
    # FILL THIS IN - END
    # -----

    # Transformer encoder does not have a last hidden layer.
    return annotations, None
```

TransformerDecoder's forward implementation

```
def forward(self, inputs, annotations, hidden_init):
    """Forward pass of the attention-based decoder RNN.

    Arguments:
        inputs: Input token indexes across a batch for all the time step. (batch_size x decoder_seq_len)
        annotations: The encoder hidden states for each step of the input sequence. (batch_size x seq_len x hidden_size)
        hidden_init: Not used in the transformer decoder

    Returns:
        output: Un-normalized scores for each token in the vocabulary, across a batch for all the decoding time steps. (batch_size x decoder_seq_len x vocab_size)
        attentions: The stacked attention weights applied to the encoder annotations (batch_size x encoder_seq_len x decoder_seq_len)
    """

    batch_size, seq_len = inputs.size()
    embed = self.embedding(inputs) # batch_size x seq_len x hidden_size

    # THIS LINE WAS ADDED AS A CORRECTION.
    embed = embed + self.positional_encodings[:seq_len]

    encoder_attention_weights_list = []
    self_attention_weights_list = []
    contexts = embed
    for i in range(self.num_layers):
        # -----
        # FILL THIS IN - START
        # -----
        new_contexts, self_attention_weights = self.self_attentions[i](contexts, annotations, annotations) # batch_size x seq_len x hidden_size
        residual_contexts = contexts + new_contexts
        new_contexts, encoder_attention_weights = self.encoder_attentions[i](residual_contexts, annotations, annotations) # batch_size x seq_len x hidden_size
        residual_contexts = residual_contexts + new_contexts
        new_contexts = self.attention_mlp[i](residual_contexts)
        contexts = residual_contexts + new_contexts

        # -----
        # FILL THIS IN - END
        # -----

        encoder_attention_weights_list.append(encoder_attention_weights)
        self_attention_weights_list.append(self_attention_weights)

    output = self.out(contexts)
    encoder_attention_weights = torch.stack(encoder_attention_weights_list)
    self_attention_weights = torch.stack(self_attention_weights_list)

    return output, (encoder_attention_weights, self_attention_weights)
```

2. Here are my results after 100 epochs

Epoch: 52	Train loss: 1.046	Val loss: 1.603	Gen: eeeway away ondinsingchgay isway orringray
Epoch: 53	Train loss: 1.034	Val loss: 1.622	Gen: eeeway away ondgiocingrgway isway odringsway
Epoch: 54	Train loss: 1.022	Val loss: 1.622	Gen: eteway away ondinsingnrgway isay oringrgay
Epoch: 55	Train loss: 1.010	Val loss: 1.669	Gen: eeeway away oncincincay-ay isway oringgray
Epoch: 56	Train loss: 1.015	Val loss: 1.655	Gen: eeeway aiday oncincionchgay isay oriningway
Epoch: 57	Train loss: 0.996	Val loss: 1.587	Gen: eteway away odcinsionggray isway oringrgway
Epoch: 58	Train loss: 0.980	Val loss: 1.602	Gen: eteway away ondinciongngway isway oringsgay
Epoch: 59	Train loss: 0.974	Val loss: 1.589	Gen: eteway away odcinsiongngway isway oringsway
Epoch: 60	Train loss: 0.972	Val loss: 1.608	Gen: eteway awrway ondinsiongngway isway orsingsway
Epoch: 61	Train loss: 0.997	Val loss: 1.664	Gen: eteway aidway onciioiongngway isway oringrgway
Epoch: 62	Train loss: 1.004	Val loss: 1.663	Gen: eteway away oddinay isway orspingway
Epoch: 63	Train loss: 0.999	Val loss: 1.599	Gen: eteway away odndiongngngway isway orsringray
Epoch: 64	Train loss: 0.969	Val loss: 1.590	Gen: eteway awsway odndinay isway orsingsway
Epoch: 65	Train loss: 0.947	Val loss: 1.596	Gen: eteway away ondinciodgrgway isway orsingsway
Epoch: 66	Train loss: 0.952	Val loss: 1.612	Gen: emeway away oncincinchngray isway orsingsway
Epoch: 67	Train loss: 0.945	Val loss: 1.563	Gen: eteway away odndiocodnggray isway odgkingway
Epoch: 68	Train loss: 0.925	Val loss: 1.576	Gen: eheway awrway odndionchnrgway isway orsingsway
Epoch: 69	Train loss: 0.912	Val loss: 1.592	Gen: eteway aidway odndiocodngngway isway orsingsway
Epoch: 70	Train loss: 0.906	Val loss: 1.566	Gen: eheway aidway ondciocongngway isway orsingsway
Epoch: 71	Train loss: 0.891	Val loss: 1.583	Gen: eheway awrway odndiocingngway isway orsingsway
Epoch: 72	Train loss: 0.886	Val loss: 1.592	Gen: eheway awrway odndiocingngway isway opringsway
Epoch: 73	Train loss: 0.879	Val loss: 1.562	Gen: eheway awrway ondingiongngway isway orringsway
Epoch: 74	Train loss: 0.874	Val loss: 1.585	Gen: eheway awrway ondincingctrway isway orsingsway
Epoch: 75	Train loss: 0.878	Val loss: 1.613	Gen: eteway irway odcincinctoday isway orpingsway
Epoch: 76	Train loss: 0.935	Val loss: 1.590	Gen: eheway aidway odndioincingway isway orkingsway
Epoch: 77	Train loss: 0.891	Val loss: 1.610	Gen: eheway irrway ondincingctrway isway orsingsway
Epoch: 78	Train loss: 0.886	Val loss: 1.624	Gen: eheway irsway ondincangctrway isway orsingsway
Epoch: 79	Train loss: 0.881	Val loss: 1.603	Gen: eheway awrway ondincinctrgway isway orringsway
Epoch: 80	Train loss: 0.862	Val loss: 1.608	Gen: eheway irrway ondincingcrgway isway orsingsway
Epoch: 81	Train loss: 0.850	Val loss: 1.597	Gen: eheway irrway ondingingctray isway opringsway
Epoch: 82	Train loss: 0.846	Val loss: 1.613	Gen: eheway aisway odcingoctingay isay opringsway
Epoch: 83	Train loss: 0.868	Val loss: 1.600	Gen: eteway aidway odndiocongngway isway orsingsway
Epoch: 84	Train loss: 0.852	Val loss: 1.607	Gen: eheway irhway ondingongctrway isway orpingsway
Epoch: 85	Train loss: 0.839	Val loss: 1.560	Gen: eheway irrway oddingonctdgway isway orringsway
Epoch: 86	Train loss: 0.825	Val loss: 1.577	Gen: eheway awrway ondingongctray isway opringsway
Epoch: 87	Train loss: 0.815	Val loss: 1.583	Gen: eheway awrway ondingodgctray isway orpingsway
Epoch: 88	Train loss: 0.807	Val loss: 1.582	Gen: eheway awrway ondiniongctrway isway orpingsway
Epoch: 89	Train loss: 0.801	Val loss: 1.581	Gen: eheway irhway ondinciongngway isway orpingsway
Epoch: 90	Train loss: 0.803	Val loss: 1.580	Gen: eheway irhway ondioiodcrgway isway orpingsway
Epoch: 91	Train loss: 0.800	Val loss: 1.566	Gen: eheway awrway ondiiocingrgway isway orpingray
Epoch: 92	Train loss: 0.795	Val loss: 1.587	Gen: eheway irrway odndioiodgrgway isway orpingray
Epoch: 93	Train loss: 0.785	Val loss: 1.597	Gen: eheway aidway ondingingctrway isway orpingsway
Epoch: 94	Train loss: 0.775	Val loss: 1.596	Gen: eheway awrway odndiongcingway isway orringray
Epoch: 95	Train loss: 0.774	Val loss: 1.589	Gen: eheway awrway ondinciongngway isway orringsway
Epoch: 96	Train loss: 0.780	Val loss: 1.606	Gen: eheway awrway ondinciongngway isway orkingray
Epoch: 97	Train loss: 0.785	Val loss: 1.619	Gen: eheway awrway ondioingcingway isway orpingsway
Epoch: 98	Train loss: 0.783	Val loss: 1.615	Gen: eheway awrway ondincingingway isway orpingsway
Epoch: 99	Train loss: 0.775	Val loss: 1.636	Gen: eheway awrway ondinciongngway isway orrkingway

source: the air conditioning is working
translated: eheway awrway ondinciongngway isway orrkingway

The model does not perform as good as I have expected, although the losses decreased a lot from the start of training, the translation is not as accurate as the additive attention model. It is also worth noting that the model reached convergence since epoch 30, so I think learning rate needs to decrease. This is also reflected by the translation on the right, which did not change that much since epoch 30.

3. After changing the init method of the transformer decoder to use only non-casual attention, I got the following results:

Epoch: 52	Train loss: 1.462	Val loss: 1.674	Gen: eway away ongningingcay ilililay oringay
Epoch: 53	Train loss: 1.459	Val loss: 1.666	Gen: eway away ongnngngngcay isway oringay
Epoch: 54	Train loss: 1.446	Val loss: 1.656	Gen: eway away ongcongway isway oringay
Epoch: 55	Train loss: 1.431	Val loss: 1.652	Gen: eway away ongnngway isililay oringingrway
Epoch: 56	Train loss: 1.423	Val loss: 1.634	Gen: eway aray onngcay isililay ongggay
Epoch: 57	Train loss: 1.410	Val loss: 1.638	Gen: ethay away onngcay isway oringay
Epoch: 58	Train loss: 1.402	Val loss: 1.633	Gen: eway away ongnngngcay isway oringingay
Epoch: 59	Train loss: 1.398	Val loss: 1.636	Gen: ethay away onngcay isililay oinggay
Epoch: 60	Train loss: 1.391	Val loss: 1.626	Gen: eway away ongnngway isililay ongggingrway
Epoch: 61	Train loss: 1.386	Val loss: 1.621	Gen: ethay awarway ongnngcay isililay ongginway
Epoch: 62	Train loss: 1.380	Val loss: 1.639	Gen: eway away ongnngongtingtway isilililway oringinway
Epoch: 63	Train loss: 1.388	Val loss: 1.623	Gen: ethay ay ongcay isay oringay
Epoch: 64	Train loss: 1.371	Val loss: 1.604	Gen: eway arway onngcay isay oringay
Epoch: 65	Train loss: 1.359	Val loss: 1.594	Gen: ethay aray ongnngngngcay isay oinggay
Epoch: 66	Train loss: 1.353	Val loss: 1.600	Gen: eway away ongnngngngngcongway isway oringingrway
Epoch: 67	Train loss: 1.345	Val loss: 1.600	Gen: ethay ararway ongnngngngngcay issilay oringingway
Epoch: 68	Train loss: 1.341	Val loss: 1.581	Gen: ethay iray ongnngngngngcay isway oringingway
Epoch: 69	Train loss: 1.324	Val loss: 1.575	Gen: ethay away ongnngngngtingcongcin issay oringingway
Epoch: 70	Train loss: 1.328	Val loss: 1.589	Gen: ethay araray ongnngngngngway isway oinggingway
Epoch: 71	Train loss: 1.323	Val loss: 1.583	Gen: ethay awaray onnglay isway oringingway
Epoch: 72	Train loss: 1.323	Val loss: 1.568	Gen: ethay away ongnngngway isway oringingway
Epoch: 73	Train loss: 1.307	Val loss: 1.584	Gen: ethay arway onngcongngngway iway oringingway
Epoch: 74	Train loss: 1.304	Val loss: 1.569	Gen: ethay arway onngngngngngway issay oringinway
Epoch: 75	Train loss: 1.290	Val loss: 1.578	Gen: ethay aray onngcongngngngway isway oringingr
Epoch: 76	Train loss: 1.289	Val loss: 1.582	Gen: ethay aray ongnngngngngway iway oringingray
Epoch: 77	Train loss: 1.286	Val loss: 1.561	Gen: ewhay away onngngiongnglay isway oringingray
Epoch: 78	Train loss: 1.283	Val loss: 1.572	Gen: ethay aray ongnngngngngcay isway oringingray
Epoch: 79	Train loss: 1.291	Val loss: 1.607	Gen: ethay aray ongnngngngngtway isway ongggggray
Epoch: 80	Train loss: 1.289	Val loss: 1.551	Gen: ethay aray ongnngngngngtway isay oringingray
Epoch: 81	Train loss: 1.280	Val loss: 1.573	Gen: eway aray ongnngngngngcay isway oringingray
Epoch: 82	Train loss: 1.280	Val loss: 1.564	Gen: ethay aray ongnngngngnggtay isway oinggingway
Epoch: 83	Train loss: 1.278	Val loss: 1.587	Gen: ethay arway ongnngngngngcay issay oringray
Epoch: 84	Train loss: 1.276	Val loss: 1.560	Gen: eway arway ongnngngngnghay isway oinggrway
Epoch: 85	Train loss: 1.256	Val loss: 1.563	Gen: ethay arway ongnngngngnghay isway oringingway
Epoch: 86	Train loss: 1.246	Val loss: 1.545	Gen: ethay arrray ongnngngngnglay isway oringingway
Epoch: 87	Train loss: 1.242	Val loss: 1.557	Gen: ethay aray ongnngngngngcay isway oringgggggggr
Epoch: 88	Train loss: 1.244	Val loss: 1.558	Gen: ethay aray ongnngngngngctingcongctin isway oinggggingay
Epoch: 89	Train loss: 1.225	Val loss: 1.560	Gen: ethay aray ongnngngngnglay isway oringggway
Epoch: 90	Train loss: 1.219	Val loss: 1.550	Gen: ethay aray ongnngngngngway isway oinggggway
Epoch: 91	Train loss: 1.226	Val loss: 1.554	Gen: ethay arway ongnngngngngctingcongctin isway ongggggr
Epoch: 92	Train loss: 1.236	Val loss: 1.542	Gen: ethay arway ongnngngngngngway isway oringingway
Epoch: 93	Train loss: 1.214	Val loss: 1.539	Gen: ethay aray ongnngngngngngcay isway oringing
Epoch: 94	Train loss: 1.211	Val loss: 1.554	Gen: ethay aray ongnngngngngngngway isway oringggway
Epoch: 95	Train loss: 1.206	Val loss: 1.545	Gen: ethay arway ongnngngngngngway isway oringr
Epoch: 96	Train loss: 1.203	Val loss: 1.545	Gen: ethay arway ongnngngngngngway isway oingr
Epoch: 97	Train loss: 1.192	Val loss: 1.531	Gen: ethay arway ongnngngngngngway isway oingr
Epoch: 98	Train loss: 1.187	Val loss: 1.543	Gen: ethay arway ongnngngngngngway isway oringing
Epoch: 99	Train loss: 1.186	Val loss: 1.535	Gen: ethay arway ongnngngngngngway isway oringing

source: the air conditioning is working
translated: ethay arway ongnngngngngngway isway oringing

Although the validation loss is slightly lower, the training loss is much higher. Also, we can see that the translation is not as good, at least in the previous one we can see traces of original words. As expected, the model reached convergence since epoch 30, and the translation is pretty much the same since epoch 30.

Part 4

1. First, I choose some samples to calculate double digit numbers that are close to each other, i.e. twelve and fourteen, these were both correct.
Second, I choose some larger number, i.e. one and hundred, they are both correct.
Third, I choose to use actual numbers, i.e. 1 and 100 to see how the model will handle this kind of input, which it was output correctly.

Fourth, I choose a mix of numbers and words, i.e. 1 and two, which “1 minus two” returned positive, a first failed case.

Fifth, I choose 2 arithmetic operations, i.e. plus and minus, which also gave correct outputs.

Lastly, I decided to use a fake number, i.e. “lala” and see how the model would perceive it.

To my surprise, it handles it pretty well. So “one minus lala” gives negative, “one plus lala” gives positive, which are some very reasonable prediction.

The following is the screenshot of my test cases:

```
[29] what_is("twelve minus fourteen")
```

☞ negative

```
[30] what_is("twelve plus fourteen")
```

☞ positive

```
[31] what_is("one plus hundred")
```

☞ positive

```
[32] what_is("one minus hundred")
```

☞ negative

```
[36] what_is("hundred minus one")
```

☞ positive

```
[37] what_is("1 plus 100")
```

☞ positive

```
[38] what_is("1 minus 14")
```

☞ negative

```
[39] what_is("1 minus two")
```

☞ positive

```
[40] what_is("two minus 3")
```

☞ negative

```
[43] what_is("three minus two minus eight")
```

☞ negative

```
[44] what_is("three minus two plus eight")
```

☞ positive

```
[45] what_is("one minus lala")
```

☞ negative

```
[46] what_is("one plus lala")
```

☞ positive

```
[47] what_is("one minus lala plus lala")
```

☞ positive

```
[49] what_is("ten minus lala")
```

☞ positive

2. For this part, I chose to change some hyperparameters, so I modified the train method to accept variables for learning rate and epochs.

I tried using 5 epochs and 3e-5 as my learning rate, the loss is as follow:

```
finttune_bert_loss_vals = train_model(model_finetune_bert, 3e-5, 5)
```

```
===== Epoch 1 / 5 =====  
Training...
```

```
Average training loss: 0.39  
Training epoch took: 0:01:12  
Running Validation...  
Accuracy: 0.94  
Validation took: 0:00:01
```

```
===== Epoch 2 / 5 =====  
Training...
```

```
Average training loss: 0.29  
Training epoch took: 0:01:12  
Running Validation...  
Accuracy: 0.97  
Validation took: 0:00:01
```

```
===== Epoch 3 / 5 =====  
Training...
```

```
Average training loss: 0.20  
Training epoch took: 0:01:12  
Running Validation...  
Accuracy: 0.98  
Validation took: 0:00:01
```

```
===== Epoch 4 / 5 =====  
Training...
```


```
Average training loss: 0.17  
Training epoch took: 0:01:12  
Running Validation...  
Accuracy: 0.98  
Validation took: 0:00:01
```

```
===== Epoch 5 / 5 =====  
Training...
```

```
Average training loss: 0.15  
Training epoch took: 0:01:13  
Running Validation...  
Accuracy: 0.98  
Validation took: 0:00:01
```

```
Training complete!
```

```
plot_loss_vals(finttune_bert_loss_vals)
```



As we can see, the loss got a bit even lower than previous model, therefore I expect that this model will perform slightly better than before.